

Lesson 7.1: Structured Query Language (SQL)

CSC430/530 – DATABASE MANAGEMENT SYSTEMS

DR. ANDREY TIMOFEEV



OUTLINE

- Introduction.
- Data definition language (DDL) commands.
 - CREATE, ALTER, DROP.
 - Domains & data types.
 - Constraints specification.
- Data manipulation language (DML) commands.
 - INSERT, DELETE, UPDATE.
 - SELECT-FROM-WHERE.

INTRODUCTION

- **Structured Query Language (SQL).**

- Most widely used relational query language.
- Serves as a standard language for storing, manipulating, and retrieving data in relational databases.

- **Data Definition Language (DDL).**

- Commands used to define and modify database schema.
- *CREATE, DROP, ALTER.*

- **Data Manipulation Language (DML).**

- Commands used to retrieve and manipulate data in a database.
- *INSERT, DELETE, UPDATE, SELECT.*

- **MySQL.**

- Open-source relational database management system.
- MySQL Workbench – IDE used for database design, development, and maintenance.
- MySQL Server – service running on a server side of client-server database management system architecture.

DDL CREATE

- **CREATE** statement allows to create a **schema** (*database*) or a **relation** (*table*).
 - **CREATE DATABASE company;** *CREATE SCHEMA is basically the same*
 - **USE company;** *allows you to refer to tables directly w/o dot operator*
 - **CREATE TABLE employee(...);** or **CREATE TABLE company.employee(...);**
↑ attributes here
- When **creating** a relation (*table*):
 - Provide a **name**;
 - Specify **attributes**, their **data types**, and **constraints**;
 - Specify **table constraints** (*optionally*).
 - Giving each constraint a name is a good database implementation practice.

DDL CREATE: DATA TYPES & DOMAIN

- Basic **data types** of attributes:
 - **Numeric** (INT, SMALLINT, FLOAT, REAL, etc.)
 - **Character-string** (CHAR, VARCHAR(n))
 - **Bit-string** (BIT(n), BIT VARYING(n))
 - **Boolean** (TRUE, FALSE, NULL)
 - **Date** (YYYY-MM-DD)
 - **Time** (HH:MM:SS)
 - **Timestamp** (DATE & TIME)
 - **Interval** (YEAR/MONTH, DAY/TIME)
- **Domain** can be explicitly created and used for multiple attributes.
 - **CREATE DOMAIN** `ssn_type` **AS** `VARCHAR(9)`;

DDL CREATE: CONSTRAINTS (1)

1. Attribute constraints:

- **NOT NULL**
 - On primary key attribute(s) (entity integrity) or any regular attribute.
- **DEFAULT** *<value>*
 - Value used if the value for an attribute is not specified.
- **CHECK**
 - Specify a certain condition.
 - **CHECK** (salary > 0);

2. Table constraints:

- **Key constraint.**
 - **PRIMARY KEY** (ssn);
- **Unique constraint.**
 - **UNIQUE** (dname);
- **Referential integrity constraint.**
 - **FOREIGN KEY** (dno) **REFERENCES** DEPARTMENT (dnumber);

3. Tuples constraints:

- **CHECK** at the end of **CREATE TABLE**
 - Applied to each tuple individually.
 - **CHECK** (dept_create_date <= mgr_start_date);

DDL CREATE: CONSTRAINTS (2)

- **Violation of referential integrity constraint is rejected** by default.
- Alternatively, **referential triggered action** can be specified:
 - ON DELETE
 - SET NULL
 - SET DEFAULT
 - CASCADE
 - ON UPDATE
 - SET NULL
 - SET DEFAULT
 - CASCADE

DDL ALTER

- **ALTER** used for several table **modifications**:
 - Adding or dropping a **column** (*attribute*).
 - Changing a **column definition**.
 - Adding or dropping **table constraints**.
- **ALTER TABLE employee**
ADD COLUMN job VARCHAR(12);
- **ALTER TABLE department**
ADD CONSTRAINT dept_mgr_fk
FOREIGN KEY (mgr_ssn) REFERENCES employee (ssn)
ON DELETE SET NULL ON UPDATE CASCADE;

DDL DROP

- **DROP** used to **drop** named schema elements.
 - Tables, domains, constraints, or schema itself.
- **Drop behavior options:**
 - CASCADE.
 - RESTRICT.
- **DROP SCHEMA company CASCADE;**
 - This removes the schema and all its elements including tables, views, constraints, etc.

DML INSERT

- **INSERT** is used to **add** one or more **row** (*tuple*) into **relation** (*table*).
 - Attribute values listed in the **same order** as specified in CREATE TABLE.
 - Rejected if any of defined **constraints** are violated.
- **INSERT INTO** **employee**
VALUES ('Richard', 'K', 'Marini', '653298653', '1962-12-30', '98 Oak Forest, Katy, TX', 'M', 37000, '123456789', 4);
- In addition, INSERT allows to assign values only for a **subset of attributes**.
 - **INSERT INTO** **employee (fname, lname, dno, ssn)**
VALUES ('Richard', 'Marini', 4, '653298653');

It won't insert if there is a violation.

DML DELETE

- DELETE is used to **remove** one or more **row** (*tuple*) from **relation** (*table*).
 - **Propagates** to other tuple(s) if **referential trigger actions** are specified.
 - Uses **WHERE** as a **condition** to select tuples to delete.
 - Missing WHERE deletes **all** tuples.

• DELETE FROM **employee**
WHERE **lname = 'Brown';**

• DELETE FROM **employee**
WHERE **ssn = '123456789';**

• DELETE FROM **employee**
WHERE **dno = 5;**

• DELETE FROM **employee;** *← deletes every Tuple in employee
(doesn't delete employee table)*

DML UPDATE

- **UPDATE** is used to **modify attribute** values of one or more selected **tuples**.
 - Uses **WHERE** as a **condition** to select tuples to update.
 - Uses **SET** to **specify** the attributes to be modified and their values.
 - Can cause **referential triggered action** if specified. ← with **ON UPDATE**
 - Updating value of **primary key** attribute will **propagate** an update in respective **foreign keys**.

• **UPDATE** **project**
 SET **plocation = 'Bellaire', dnum = 5**
 WHERE **pnumber = 10;**

• **UPDATE** **employee**
 SET **salary = salary * 1.1**
 WHERE **dno = 5;**

} semantically, everyone in dept 5
 is getting a raise.

DML SELECT (1)

- **SELECT** is used to **retrieve** specific data form the database.
- Basic form of SELECT statement (*select-from-where*):
 - **SELECT** *<attribute list>*
 FROM *<table list>*
 WHERE *<condition>;*
 - <attribute list> - **attribute names** whose values are to be retrieved.
 - <table list> - **relation names** required to process the query.
 - <condition> - **Boolean expression** that identifies the tuples to be retrieved by the query.
- Select birth date and address of employee John B Smith.
SELECT **bdate, address**
FROM **employee**
WHERE **fname = 'John' AND minit = 'B' AND lname = 'Smith';**

DML SELECT (2)

- **SELECT-PROJECT-JOIN** query:

- Select first name, last name and address of all employees who work for Research department.

```
SELECT    fname, lname, address
FROM      employee, department
WHERE     dname = 'Research' AND dnumber = dno;
```

- Select last name, address, and birth date of employees who manage departments with projects located in Stafford.

```
SELECT    pnumber, dnum, lname, address, bdate
FROM      project, department, employee
WHERE     dnum = dnumber AND mgr_ssn = ssn AND plocation = 'Stafford';
```

DML SELECT (3)

- **Prefixing** is used when referencing two (or more) attributes with the same name in different relations.
 - *employee.name* and *department.name*
- **Aliasing of relations** (*tuple variables*) is used to rename a relation with an abbreviation.
 - Useful when referring to the same relation twice.
 - **Example:** Select first name and last name of employees and their supervisors.
 - **SELECT** **e.fname, e.lname, s.fname, s.lname**
 FROM **employee e, employee s**
 WHERE **e.super_ssn = s.ssn;**
- **Aliasing of attributes** can be done in SELECT part of the query.
 - **SELECT** **fname AS fn, lname AS ln, bdate AS bd**
- **Missing WHERE** selects all tuples (single relation) or does CROSS PRODUCT (multiple relations).
- **Asterisk *** used to select all the attributes (no projection / projection on all attributes).