

W22

CSC 430/530 DBMS/DT

Lab 2 - Examples

8

Basic SQL Query

SELECT	[DISTINCT] <attribute list>
FROM	<table list>
WHERE	<condition>;

- table-list or relation-list: List of relation names
- Attribute list or target-list: List of attributes of tables in *relation-list*
- qualification: Comparisons combined using AND, OR and NOT.
- DISTINCT: optional keyword indicating that the answer should not contain duplicates.

9

Query Semantics

1. FROM : compute *cross product* of tables.
2. WHERE : Check conditions, discard tuples that fail.
3. SELECT : Delete unwanted fields.
4. DISTINCT (*optional*) : eliminate duplicate rows.

Note: Probably the least efficient way to compute a query!
Query optimizer will find more efficient ways to get
 the *same answer*.

10

Example of Conceptual Evaluation (1)

SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=103;

(1) Compute the cross-product of
 relation-list.

Sailors					Reserves		
<u>sid</u>	sname	rating	age		<u>sid</u>	<u>bid</u>	day
22	dustin	7	45.0	x	22	101	10/10/96
31	lubber	8	55.5		58	103	11/12/96
58	rusty	10	35.0				

11

Example of Conceptual Evaluation (1)

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=103;
```

(1) Compute the cross-product of relation-list.

Sailors					Reserves		
<u>sid</u>	sname	rating	age		<u>sid</u>	<u>bid</u>	day
22	dustin	7	45.0	x	22	101	10/10/96
31	lubber	8	55.5		58	103	11/12/96
58	rusty	10	35.0				

12

Example of Conceptual Evaluation (2)

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=103;
```

(2) Discard tuples if they fail qualifications.

Sailors X Reserves						
S.sid	sname	rating	age	R.sid	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

13

Example of Conceptual Evaluation (2)

```
SELECT S.sname
FROM Sailors S, Reserves R
```

```
WHERE S.sid=R.sid AND R.bid=103;
```

(2) Discard tuples if they fail qualifications.

Sailors X Reserves

S.sid	sname	rating	age	R.sid	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

14

Example of Conceptual Evaluation (2)

```
SELECT S.sname
FROM Sailors S, Reserves R
```

```
WHERE S.sid=R.sid AND R.bid=103;
```

(2) Discard tuples if they fail qualifications.

Sailors X Reserves

S.sid	sname	rating	age	R.sid	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

15

Example of Conceptual Evaluation (3)

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=103;
```

(3) Delete attribute columns that are not in target-list.

							sname
Sailors X Reserves							rusty
(sid)	sname	rating	age	(sid)	bid	day	
22	dustin	7	45.0	22	101	10/10/96	
22	dustin	7	45.0	58	103	11/12/96	
31	lubber	8	55.5	22	101	10/10/96	
31	lubber	8	55.5	58	103	11/12/96	
58	rusty	10	35.0	22	101	10/10/96	
58	rusty	10	35.0	58	103	11/12/96	

16

Aliasing

Consider the following SALESREPS relation

Empl_num	name	age	Rep_office	manager
105	Bill	37	13	104
104	Bob	33	12	106
106	Sam	52	11	NULL

Query: Determine the name of Bob's manager?

HINT: Use the ρ operator from Relational Algebra

17

Aliasing in SQL

```
SELECT s2.name
FROM   SALESREPS s1, SALESREPS s2
WHERE  s1.name='Bob' AND
       s1.manager=s2.empl_num;
```

- Aliases must be used here.
- The row referenced by s1 is intended to be Bob (employee role)...while s2 will be his manager (role).
- Keep in mind SQL executes its queries: first FROM, then WHERE, then SELECT

18

Important naming of variables

- Needed when ambiguity could arise.
 - e.g., same table used multiple times in FROM (“self-join”)

```
SELECT  x.sname, x.age, y.sname, y.age
FROM    Sailors x, Sailors y
WHERE   x.age > y.age
```

Sailors x

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Sailors y

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

21

Arithmetic Expressions

```
SELECT S.age, S.age-5 AS age1, 2*S.age AS age2
FROM   Sailors S
WHERE  S.sname = 'dustin'
```

```
SELECT S1.sname AS name1, S2.sname AS name2
FROM   Sailors S1, Sailors S2
WHERE  2*S1.rating = S2.rating - 1
```

22

String Comparisons

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sname LIKE 'B_%B'
```

'_' stands for any one character and '%' stands for 0 or more arbitrary characters.

23

Find sid's of sailors who've reserved a red or a green boat

```
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND
        (B.color='red' OR B.color='green')
```

... Or:

```
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND B.color='red'
UNION
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND B.color='green'
```



24

Find sid's of sailors who've reserved red and green boats

```
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid
        AND (B.color='red' AND
              B.color='green')
```



25

Find sid's of sailors who've reserved red and green boats

```
SELECT S.sid
FROM   Sailors S, Boats B, Reserves R
WHERE  S.sid=R.sid
        AND R.bid=B.bid
        AND B.color='red'
INTERSECT
SELECT S.sid
FROM   Sailors S, Boats B, Reserves R
WHERE  S.sid=R.sid
        AND R.bid=B.bid
        AND B.color='green'
```



26

Find sid's of sailors who've reserved red and green boats

- Could use joins:

```
SELECT R1.sid
FROM   Boats B1, Reserves R1, Boats B2, Reserves R2
WHERE  R1.sid=R2.sid
        AND R1.bid=B1.bid
        AND R2.bid=B2.bid
        AND (B1.color='red' AND B2.color='green')
```



27

Find sid's of sailors who have not reserved a boat

```
SELECT S.sid
FROM   Sailors S
EXCEPT
SELECT S.sid
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid
```



28

Nested Queries: IN

Names of sailors who've reserved boat #103:

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid IN (SELECT  R.sid
                  FROM    Reserves R
                  WHERE    R.bid=103)
```

29

Nested Queries: NOT IN

Names of sailors who've not reserved boat #103:

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid NOT IN (SELECT  R.sid
                     FROM      Reserves R
                     WHERE     R.bid=103)
```

30

Motivation for Subqueries

- Find the name of the professors who teach “CS 430.”

```
SELECT Name
FROM    Professors, Teach
WHERE   (PID = ProfessorPID) AND (Number = '430')
        AND (DeptName = 'CS');
```

- Do we need to take the natural join of two big relations just to get a relation with few tuples?
- Can we rewrite the query without using a join?

32

Nesting

- A query can be put inside another query
- Most commonly in the WHERE clause
- Sometimes in the FROM clause (depending on the DBMS)
- This subquery is executed first (if possible)

33

Subquery Example

- Find the name of the professor who teaches “CS 430.”

```
SELECT Name
FROM Professors
WHERE PID =
    (SELECT ProfessorPID
     FROM Teach
     WHERE (Number = 430) AND (DeptName = 'CS')
    );
```

- When using `=`, the sub-query must return a single tuple

34

Conditions Involving Relations

- SQL includes operators that can be applied to produce a Boolean result.
- These operators are useful while working with sub-queries.
- Let **R** be a relation and **t** be a tuple (with the same set of attributes as that of R).
 1. **t EXISTS R** is true if and only if R contains at least one tuple.
 2. **t IN R** is true if and only if t equals a tuple in R.
 3. **t > ALL R** is true if and only if R is *unary* (has one attribute) and t is greater than **every** value in R.
 4. **t > ANY R** (which is *unary*) is true if and only if t is greater than **at least** one value in R.
- We can use **NOT** to negate EXISTS, ALL, and ANY.

35

Subqueries Using Conditions

- Find the departments of the courses taken by the student with name 'Suri'.

```
SELECT DeptName
FROM Take
WHERE StudentPID IN
    ( SELECT PID
      FROM Students
      WHERE (Name = 'Suri')
    );
```

36

Correlated vs Uncorrelated

- The previous sub-queries did not depend on anything **outside the sub-query**
 - ...and thus need to be executed just once.
 - These are called uncorrelated.


A correlated sub-query depends on data from the outer query ... and thus must be executed for each row of the outer table(s)

37

Nested Queries with Correlation

Names of sailors who've reserved boat #103:

```
SELECT  S.sname
FROM    Sailors S
WHERE   EXISTS
        (SELECT *
         FROM Reserves R
         WHERE R.bid=103 AND S.sid=R.sid)
```



- Sub-query must be **recomputed** for each Sailors tuple.
 - Think of sub-query as a function call that runs a query!
- The same holds for NOT EXISTS as well.

38

Correlated Subqueries

- Find course names that have been used for two or more courses.

```
SELECT CourseName
FROM Courses AS First
WHERE CourseName IN
    (SELECT CourseName
     FROM Courses
     WHERE (Number <> First.Number)
     AND (DeptName <> First.DeptName)
    );
```

39

Evaluating Correlated Subqueries

```
SELECT CourseName
FROM Courses AS First
WHERE CourseName IN
    (SELECT CourseName
     FROM Courses
     WHERE (Number <> First.Number)
     AND (DeptName <> First.DeptName)
    );
```

- Evaluate query by looping over tuples of First, and for each tuple evaluate the subquery.
- Scoping rules:** an attribute in a subquery belongs to one of the tuple variables in that subquery's FROM clause, or to the immediately surrounding subquery, and so on.

40

UNIQUE

*Names of sailors who've reserved boat #103
exactly once:*

```
SELECT  S.sname
FROM    Sailors S
WHERE   UNIQUE
        (SELECT  *
         FROM    Reserves R
         WHERE   R.bid=103 AND
                 S.sid=R.sid)
```

42

A Tough One

Find sailors who've reserved all boats.

```
SELECT  S.sname      Sailors S such that ...
FROM    Sailors S
WHERE   NOT EXISTS (SELECT  B.bid
                   FROM    Boats B there is no boat B without ...
                   WHERE   NOT EXISTS (SELECT  R.bid
                                       FROM    Reserves R
                                       WHERE   R.bid=B.bid
                                       AND R.sid=S.sid))
a Reserves tuple showing S reserved B
```

44

Summary

- Relational model has well-defined query semantics
- SQL provides functionality close to basic relational model
(*some differences in duplicate handling, null values, set operators, ...*)
- Typically, many ways to write a query
 - DBMS figures out a fast way to execute a query, regardless of how it is written.