

CSC 430/530 Database Management
Systems / Database Theory

PHYSICAL QUERY PLAN

LECTURE 20

28

Query Processing

- Efficient Query Processing crucial for good or even effective operations of a database
- Query Processing depends on a variety of factors, not everything under the control of DBMS
- Insufficient or incorrect information can result in vastly ineffective plans
- Query Cataloging

29


Physical Query Plan Operators

- Basic set of operators that define the language of physical query execution plans
- Comprises the set of relational operators in addition to some more required operators

30

Physical Query Plan Operators

- Table scan operator:
 - Scan and return an entire relation R
 - Scan and return only those tuples of R that satisfy a given predicate
- Table scan: reading all blocks of a sorted data file in sequence
- Index scan: Use an index to read all blocks of a data file in sequence



31

Sorting while Scanning

- Sort-scan operator sorts a relation R while scanning it into memory
 - If sorting is done on an indexed key attributed, no need to do anything; scanning will read R in sorted order
 - If R is small enough to fit in main memory, perform sorting after scanning
 - Else, use external sort-merge techniques to implement sort-scan

32

Iterators

- Iterators are physical-query-plan operators that comprise of three stages:
 - `Open()` where an iterable object (such as a relation is opened)
 - `GetNext()` returns the next element of the object
 - `Close()` closes control on the object

33

Iterators (Example)

- Example of a “Table scan” iterator:

```

Open() {
    b ← first block of R;
    t ← first tuple of b;
}

GetNext () {
    if (t is beyond the last tuple in b) {
        increment b;
        if (b is beyond last block) RETURN NoMoreData;
        else t ← first tuple of b;
    }
}

```

34

Iterators (Example)

```

        oldt ← t;
        increment t;
        RETURN oldt;
    }

    Close () {
    }
}

```

35

Iterators (Example 2)

- Computing Bag union $R+S$ using iterators over iterators

```
Open (){
  R.Open();
  CurRel ← R;
}
```

```
GetNext (){
  If (CurRel = R) {
    t ← R.GetNext();
    if (t != NoMoreData)
      return t;
```

36

Iterators (Example 2)

```
    else { /* R is exhausted */
      S.Open();
      CurRel ← S;
    }
  } ELSE { /* if (CurRel = R ) */
    Return S.GetNext();
  }

  Close () { R.Close(); S.Close(); }
```

37

Database Access Algorithms

- Algorithms for database access can be broadly divided in the following categories:
 - Sorting-based methods
 - Hash-based methods
 - Index-based methods

38

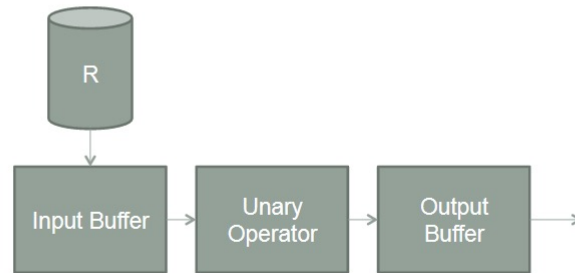
Physical Query Plan Operator Types

- Tuple-at-a-time unary operators:
 - Can read only one block at a time and required to work with only one tuple
- Full-relation unary operators:
 - Requires knowledge of all or most of the relation. Read into memory for small relations in one-pass algorithms
- Full-relation binary operators:
 - Same as above, but on two relations.

39

Tuple-at-a-time Unary Operations

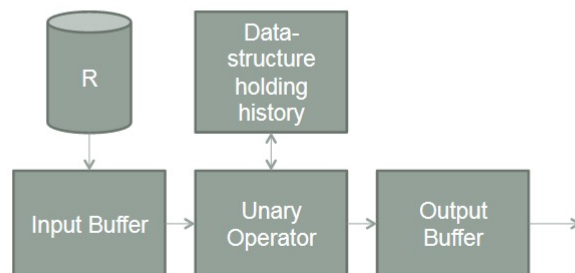
- Examples: $\sigma(R)$, $\pi(R)$ etc...
- A strategy for a **one-pass algorithm**:



40

Relation-at-a-time Unary Operations

- Examples: UNIQUE, GROUP BY, etc...
- A strategy for a **one-pass algorithm**:



41

Relation at a time Binary Operators

- One-pass strategies for binary relation-at-a-time operators vary between different operators.
- Almost all of them require at least one of the relation to be completely stored in memory

42

Strategies

- Set Union $R \cup S$
- Assuming R is the bigger relation:
 - Read S into memory completely and make it accessible through an in-memory index structure
 - Output all tuples of S while reading
 - For each tuple of R , search if it already exists in S , and output if not.

43

Strategies

- Set Intersection $R \cap S$
- Assuming R is the bigger relation:
 - Read S into memory completely and make it accessible through an in-memory index structure
 - For each tuple of R, search if it already exists in S, and output if not.

44

Strategies

- Set Difference $R - S$
- Assuming R is the bigger relation:
 - Read S into memory completely and make it accessible through an in-memory index structure
 - For each tuple of R, search if it already exists in S.
 - If tuple exists in S, then ignore; else output the tuple.

45

Strategies

- Set Difference S-R
- Assuming R is the bigger relation:
 - Read S into memory completely and make it accessible through an in-memory index structure
 - For each tuple of R, search if it already exists in S, and delete it from S if it exists.
 - Output all remaining tuples of S.

46

Strategies

- Cross Product $R \times S$
- Assuming R is the bigger relation:
 - Read S into memory completely and store it in a buffer. No special data structure required
 - For each tuple of R, combine it with each tuple of S and output the result.

47

Strategies

- Natural Join $\bar{R} \bowtie \bar{S}$. Assume $R(X,Y)$ and $S(Y,Z)$
- Assuming R is the bigger relation:
 - Read S into memory completely and store it in a balanced tree index structure or a hash table.
 - For each tuple of R , search S to see if a matching tuple exists.
 - Output if matching tuple found.

48

One-pass algorithms

- One-pass algorithms are applicable only when one of the relation fits completely into memory.
- In addition, there is enough memory to store at least one block of the other relation
- Hence, if M memory buffers are available, then one of the relations should have a maximum size of $M-1$.

49

One-pass algorithms

- One-pass algorithms rely on correctly estimating relation sizes and allocating memory buffers.
- If too many buffers are allocated, there is a possibility of **thrashing**
- If too few buffers are allocated, then one-pass algorithms may not run

50

Summary

- Stages in Query Processing
- Logical Query Plan and Physical Query Plan
- Intermediate Query Language
- Physical Query Plan Language constructs
- One-pass algorithm for unary and binary operators.

51

Multi-pass Algorithms

- Used when entire relations cannot be read into memory
- Requires alternate computation and retrieval of intermediate results
- Many multi-pass algorithms are generalizations of their corresponding two pass algorithms

52

Basic idea: Two-pass Algorithms based on Sorting

- Suppose relation R is too big to fit in memory which can accommodate only M blocks of data.
- The “sorting-based” 2-pass algorithms have the following basic structure:
 1. Read M blocks of records into memory and sort them
 2. Write them back to disk
 3. Continue steps 1 and 2 until R is exhausted
 4. Use a variety of “query-merge” techniques to extract relevant results from all the sorted M-blocks on disk

53

Example of two-phase algorithm:

- Query: Duplicate elimination using sorting
- Phase 1:
 - Let relation R , in which duplicates must be eliminated, be too big to fit in memory.
 - **Step A:** Read $M \ll R$ blocks into memory and sort them.
 - Store the sorted set of M blocks back on disk
 - Return to step A till R is exhausted.

54

Example of two-phase algorithm:

- Phase 2:
 - Take one block from each sorted sub-list on disk and eliminate duplicate tuples
 - **More specifically:** Take the first element p of the first block and move all other blocks to go beyond p .
 - Since blocks are sorted, the “merge-elimination” take $O(M \times n \times b)$, where M is the number of blocks in a sorted block set, n is the number of block sets, and b is the number of tuples in a block.

55

Duplicate elimination using sorting

	In Memory	Waiting on Disk
1	1, 2	2, 2 2, 5
2	2, 3	4, 4 4, 5
3	1, 1	2, 3 5

Output 1 from the first block and move all blocks beyond "1"

	In Memory	Waiting on Disk
1	2	2, 2 2, 5
2	2, 3	4, 4 4, 5
3	2, 3	5

56

Duplicate elimination using sorting

	In Memory	Waiting on Disk
1	2	2, 2 2, 5
2	2, 3	4, 4 4, 5
3	2, 3	5

Output 2 from the first block and move all blocks beyond "2"

	In Memory	Waiting on Disk
1	5	
2	3	4, 4 4, 5
3	3	5

57

Example of two phase algorithm: (2)

Set Union using sorting is simply duplicate elimination from two sets.

For computing $R \cup S$:

1. Read blocks from R and S, sort them and store them on disk.
2. Repeat step 1 until R and S are exhausted.
3. Use duplicate elimination over the set of all these blocks outlined in the previous example.

58

Example of two phase algorithm: (3)

Set intersection using sort:

1. Given two relations R and S, read them in terms of blocks and store the blocks on disk in sorted order
2. Take each block of R stored on disk and perform the following:
 1. Read the first tuple t from R's block
 2. Move all blocks of S beyond t (i.e., travers until t in S)
 3. If t existed in any block of S, output t, else ignore and move onto the next tuple in R.
3. (Note: Set intersection assumes R and S are sets; i.e. no duplicates)

59

Example of two phase algorithm: (4)

Natural Join Using Sort

In order to compute $R(X,Y) \bowtie S(Y,Z)$ using sorting:

1. Read blocks of R and sort them on the Y attribute and store them back on disk
2. Read blocks of S and sort them on the Y attribute and store them back on disk
3. For each value y at the top of the first block of R:
 1. If y occurs in S, then read all blocks of R beyond y and store them separately

60

Summary

- Multi pass algorithm
 - We know the two-phase algorithm
 - Duplicate elimination is important
 - Review basic operators

61